


Big Data Infrastructure and Technologies for Data Analytics (BDIT4DA)

Practice Guidelines

Hadoop cluster Installation:
Cloudera Quickstart on VirtualBox




Outline

- Hadoop cluster VMs
 - Cloudera Qyquickstart 5.13 – 5.3 GB (discontinued but version 5.13 available)
 - Hortonworks Sandbox – 10 GB
 - Oracle Big Data Lite 4.11 – 22 GB
 - Bitnami Hadoop Stack – 1.5 GB (only CLI and MapReduce)
- Oracle VirtualBox and pre-requisite configuration
- Installing Cloudera Quickstart Hadoop cluster on VirtualBox


Preparation

- Download Cloudera Quickstart 5.13 from below location (not available at Cloudera website since 2020)
 - <https://surfdrive.surf.nl/files/index.php/s/2OvUuw5chV42Zwz>
 - Directory “vm-software”
- Configure VirtualBox Host-Only Network Adapter




Cloudera Quickstart VM for VirtualBox

- QuickStart virtual machines (VMs) provide sufficient functionality to try CDH, Cloudera Manager, Impala, and Cloudera Search
 - Minimum required RAM on host machine is 4GB, recommended 8GB
- Cloudera Manager is installed in the VM but is turned off by default.
 - Cloudera strongly recommends that if you use Cloudera Manager, you configure the VM with a minimum of 8 GB RAM and two virtual CPU cores.



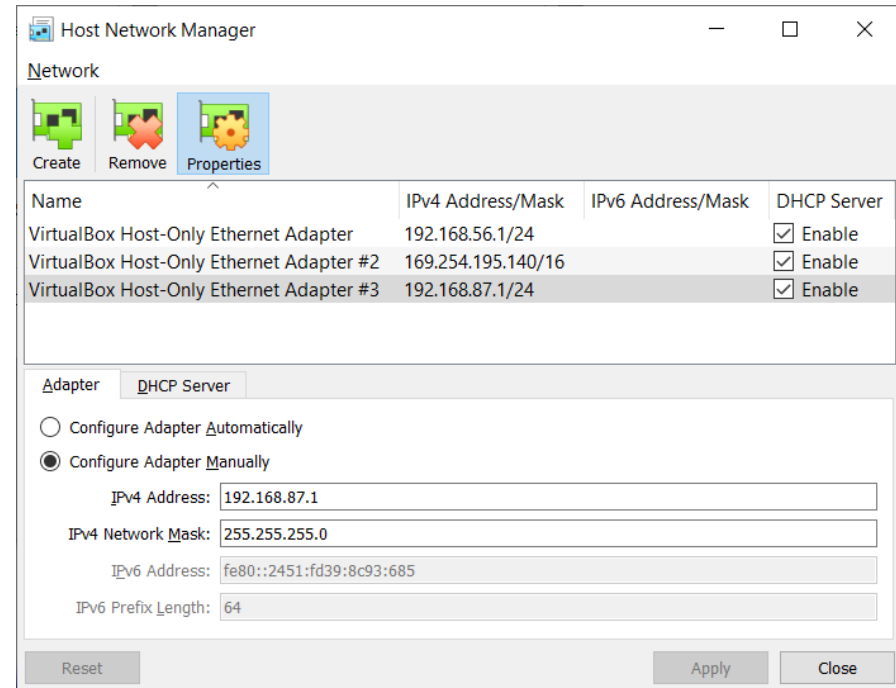
VirtualBox Configuration – Host requirements

- Pre-requisite and requirement to host
- Hosting Cloudera Quickstart 5.13 VM require
 - Single node minimum RAM 4GB, recommended 8GB
 - Single node + Cloudera Manager minimum 8GB
- Host machine HardDrive space:
 - Quickstart 5.13 VM for VirtualBox – 5.3 GB
 - VM image installed on the host machine – 9-12 GB
 - Plus additionally for a snapshot 1.2-2.5 GB
- VirtualBox Installation
 - **Note: For Windows, you can configure VirtualBox VMs location on drive different from C: (which may be have limited space)**



Accessing your VirtualBox Guest from your Host OS: Adding Host-Only network interface

- You need to do is to setup a VirtualBox Host-Only Network and get our guest connected to it.
 - If your guest machine is running, shut it down first.
- New VirtualBox versions
 - Click on File->Host Network Manager in the VirtualBox menu-bar.
 - Use Create/Properties Network Adapter
- Older VirtualBox versions
 - Click on File->Preferences
 - Select the Network option from the side menu and add network adapter
- Configure Adapter Automatically or Manually for better control
 - The default options for the newly-created Host-only network should be fine.
 - If not, you can add the data manually, by clicking on the Edit button in the DHCP Server tab.
 - Save all the settings in Preferences.
- Now open up the settings of your Guest machine and navigate to the Network option from the side menu and click on the Adapter 2 tab.
 - Don't forget to check the Enable Network Adapter option.
 - Save these settings and boot into your Guest machine.
 - After logging in, type `ifconfig`. Note the new IP, it should be under a new interface like `eth0` or `vboxnet0`. Now you can use this IP to SSH, view the webpages on your machine's Apache Server, etc..



Importing VM: Select+Click on .ova File

cloudera-quickstart-vm-5.13.0-0-virtualbox

File Explorer ribbon: Clipboard, Organize, New, Open, Select

File Explorer address bar: virtualbox-vm > cloudera-quickstart-vm-5.13.0-0-virtualbox

Name	Date modified	Type	Size
cloudera-quickstart-vm-5.13.0-0-virtualb...	03-May-20 2:13 PM	Open Virtualizatio...	15 KB
cloudera-quickstart-vm-5.13.0-0-virtualb...	03-May-20 2:13 PM	Virtual Machine Di...	5,810,349 ..

Oracle VM VirtualBox Manager

File Machine Help

Tools

Import Virtual Appliance

Appliance settings

These are the virtual machines contain... machines. You can change many of th... check boxes below.

Importing Appliance ...: Importing appliance 'D:\VirtualBoxVMs\vm_cloudera-quickstar...'

Importing virtual disk image 'cloudera-quickstart-vm-5.13.0-0-virtualbox-disk1.vmdk' ... (2/3)

32% 5 minutes remaining

Virtual System 1

Name	Value
Name	cloudera-quickstart-vm-5.13.0-0-virtualbox
Guest OS Type	Red Hat (64-bit)
CPU	2
RAM	8192 MB
DVD	<input checked="" type="checkbox"/>
Network Adapter	<input checked="" type="checkbox"/> Intel PRO/1000 MT Desktop (82540EM)
Storage Controller (IDE)	PIIX4


Machine Base Folder: D:\VirtualBoxVMs

MAC Address Policy: Include only NAT network adapter MAC addresses

Additional Options: ☒ Import hard drives as VDI

Appliance is not signed

Restore Defaults Import Cancel



VirtualBox Host-Only Network Configuration

Oracle VM VirtualBox Manager

File Machine Help

Tools

New Settings Discard Start

bitnami-lampstack (Snapshot 1) Saved

cloudera-quickstart-v... Saved

cloudera-quickstart-vm-5.13.0... Saved

Hortonworks Sandbox HDP 3.0 Saved

General

Name: cloudera-quickstart-vm-5.13.0-0-virtualbox

Operating System: Red Hat (64-bit)

System

Base Memory: 8192 MB

Processors: 2

Boot Order: Hard Disk, Optical, Floppy

Acceleration: VT-x/AMD-V, KVM Paravirtualization

Display

Video Memory: 16 MB

Graphics Controller: VBoxSVGA

Remote Desktop Server: Disabled

Recording: Disabled

Storage

Controller: IDE Controller

IDE Primary Master: cloudera-quickstart-vm-5.13.0-0-virtualbox

IDE Secondary Master: [Optical]

Host Network Manager

Network

Create Remove Properties

Name	IPv4 Address/Mask	IPv6 Address/Mask	DHCP Server
VirtualBox Host-Only Ethernet Adapter	192.168.56.1/24		<input checked="" type="checkbox"/> Enable
VirtualBox Host-Only Ethernet Adapter #2	169.254.195.140/16		<input checked="" type="checkbox"/> Enable
VirtualBox Host-Only Ethernet Adapter #3	192.168.87.1/24		<input checked="" type="checkbox"/> Enable

Adapter DHCP Server

☐ Configure Adapter Automatically

☒ Configure Adapter Manually

IPv4 Address: 192.168.87.1

IPv4 Network Mask: 255.255.255.0


IPv6 Address: fe80::2451:fd39:8c93:685

IPv6 Prefix Length: 64

Reset Apply Close

- File > Host Network Manager
- > Create Host-Only Network

- Configure Host-Only Adapters with addresses of your local network addresses
- Beneficially do it e.g. with address range of home and work network



Guest VM Network Configuration

Oracle VM VirtualBox Manager

File Machine Help

Tools

bitnami-lampstack (Snapshot 1) Saved

cloudera-quickstart-v... Saved

cloudera-quickstart-vm-5.13.0... Saved

Hortonworks Sandbox HDP 3.0 Saved

General

Name: cloudera-quickstart-vm-5.13.0-0-virtualbox

Operating System: Red Hat (64-bit)

System

Base Memory: 8192 MB

Processors: 2

Boot Order: Hard Disk, Optical

Acceleration: VT-x/AMD-V, Nested VT-x/AMD-V, KVM Paravirtualization

Display

Video Memory: 16 MB

Graphics Controller: VBoxVGA

Remote Desktop Server: Disabled

Recording: Disabled

Storage

Controller: IDE Controller

IDE Primary Master: cloudera-quickstart-vm-5.13.0-0-virtualbox

IDE Secondary Master: [Optical Drive]

cloudera-quickstart-vm-5.13.0-0-virtualbox - Settings

Network

Adapter 1 Adapter 2 Adapter 3 Adapter 4

☒ Enable Network Adapter

Attached to: Host-only Adapter

Name: VirtualBox Host-Only Ethernet Adapter #2

Advanced

Invalid settings detected

OK Cancel

- Settings > Network
- Adapter 1 > Host-Only Network
- Adapter 2 > NAT

- Configure Host-Only Adapters with addresses of your local network addresses
- Beneficially do it e.g. with address range of home and work network

Add Shared Folder: e.g. Exercise_Files_all

The screenshot shows the Oracle VM VirtualBox Manager interface. On the left, a list of VMs includes 'bitnami-lampstack (Snapshot 1)', 'cloudera-quickstart-v...', 'cloudera-quickstart-vm-5.13.0...', and 'Hortonworks Sandbox HDP 3.0'. The 'cloudera-quickstart-vm-5.13.0...' VM is selected. The 'Settings' button is highlighted with a blue arrow pointing to the 'Shared Folders' tab in the settings window.

cloudera-quickstart-vm-5.13.0-0-virtualbox - Settings

Shared Folders

Name	Path	Access	Auto Mount	At
Ex...ll	D:\...EV\Exercise_Files_all	Full	Yes	exercises

Add Share

Folder Path: D:\...EV\Exercise_Files_all

Folder Name: Exercise_Files_all

☐ Read-only

☒ Auto-mount

Mount point: Exercises

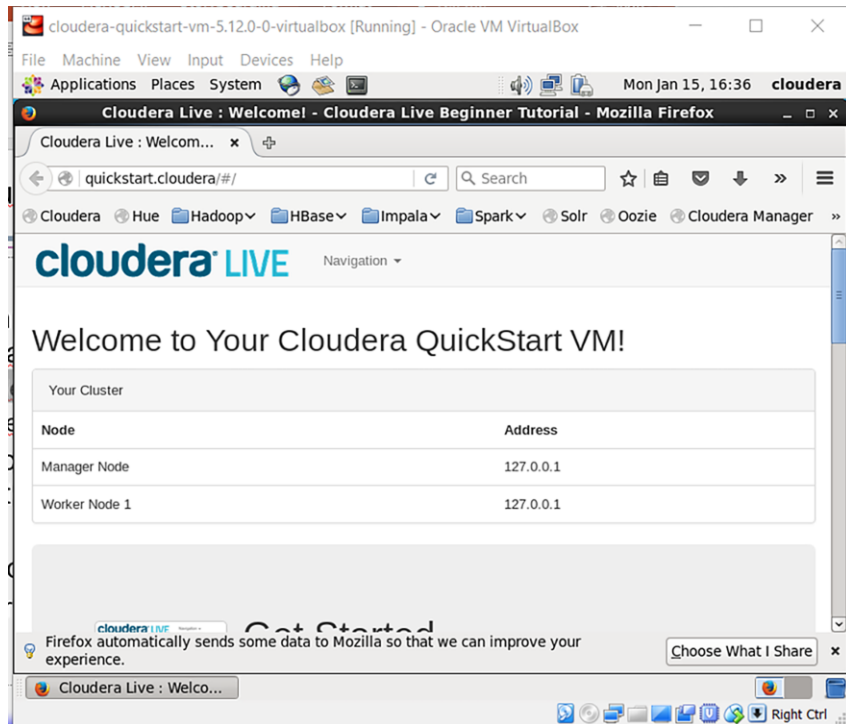
OK Cancel

Invalid settings detected

- Settings > Shared Folders
- Add Share > Folder Path > other
- Select folder + Auto-mount
- Mount point


- Advice: Place all your working files in one folder on your drive

Cloudera Quickstart VM for VirtualBox



Accounts

- Once you launch the VM, you are automatically logged in as the cloudera user. The account details are:
 - username: cloudera
 - password: cloudera
- The cloudera account has sudo privileges in the VM. The root account password is cloudera.
- The root MySQL password (and the password for other MySQL user accounts) is also cloudera.
- Hue and Cloudera Manager use the same credentials.



Find IP address of your VM

- Debian
 - \$ ip address show
- Ubuntu
 - \$ ipconfig
- Example:
 - 192.168.102.154 – LAN
 - 192.168.163.3 – VirtualBox Host only
- Use for configuring your SSH or SCP client

Sun Apr 26, 5:32 PM cloudera


Cloudera Live Beginner Tutorial - Mozilla Firefox

start VM!

	Address
File Machine View Inp	192.168.102.164
Activities Terminal	192.168.102.164

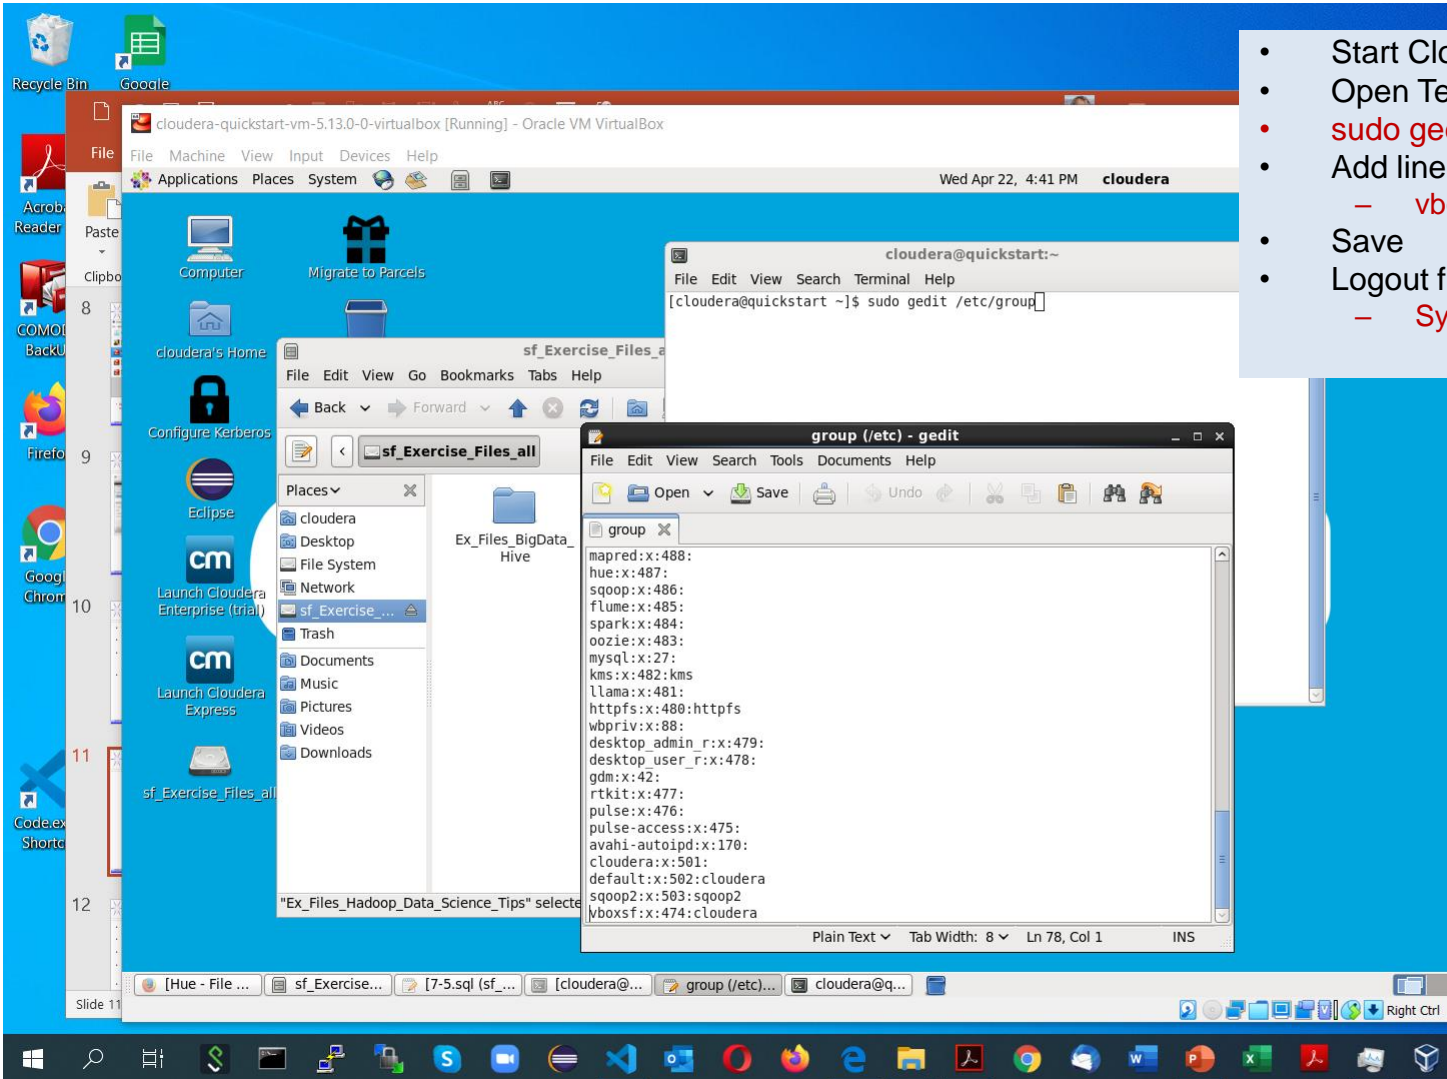
```
File Edit View Search Terminal Help
[cloudera@quickstart ~]$ ip address show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
   inet 127.0.0.1/8 scope host lo
       link/ether 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:b2:30:58 brd ff:ff:ff:ff:ff:ff
   inet 192.168.102.164/24 brd 192.168.102.255 scope global eth0
       link/ether 08:00:27:fb:f6:a9 brd ff:ff:ff:ff:ff:ff
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:fb:f6:a9 brd ff:ff:ff:ff:ff:ff
   inet 10.0.3.15/24 brd 10.0.3.255 scope global eth1
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:84:a2:84 brd ff:ff:ff:ff:ff:ff
   inet 192.168.56.102/24 brd 192.168.56.255 scope global eth2
[cloudera@quickstart ~]$
```


```
File Edit View Search Terminal Help
mysql> exit
Bye
bitnami@debian:~$
Enter password:
ERROR 1045 (28000)
(ES)
bitnami@debian:~$
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
   link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
   inet 127.0.0.1/8 scope host lo
       link/ether 00:00:00:00:00:00 brd 00:00:00:00:00:00
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:a9:a6:22 brd ff:ff:ff:ff:ff:ff
   inet 192.168.102.154/24 brd 192.168.102.255 scope global eth0
       link/ether 08:00:27:fb:f6:a9 brd ff:ff:ff:ff:ff:ff
       valid_lft forever preferred_lft forever
   inet6 fe80::a03:27ff:faa9:a622/64 scope link
       valid_lft forever preferred_lft forever
3: eth1: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:ba:93:bc brd ff:ff:ff:ff:ff:ff
   inet 10.0.3.15/24 brd 10.0.3.255 scope global dynamic eth1
       link/ether 840805cc preferred_lft 840805cc
       valid_lft 493sec preferred_lft 493sec
   inet6 fe80::a03:27ff:faa9:a622/64 scope link
       link/ether 08:00:27:84:a2:84 brd ff:ff:ff:ff:ff:ff
       valid_lft forever preferred_lft forever
4: eth2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP qlen 1000
   link/ether 08:00:27:6a:8f:38 brd ff:ff:ff:ff:ff:ff
   inet 192.168.87.3/24 brd 192.168.87.255 scope global dynamic eth2
       link/ether 840805cc preferred_lft 840805cc
       valid_lft 493sec preferred_lft 493sec
   inet6 fe80::a03:27ff:faa9:a622/64 scope link
       link/ether 08:00:27:fb:f6:a9 brd ff:ff:ff:ff:ff:ff
       valid_lft forever preferred_lft forever
bitnami@debian:~$
```



Add Permission for User "cloudera"


- Start Cloudera
- Open Terminal
- `sudo gedit /etc/group`
- Add line
 - `vboxsf:x:474:cloudera`
- Save
- Logout from the system & Login
 - `System > Logout cloudera`





Working with Cloudera Quickstart

- Mount your local/host directory with exercises to Cloudera HDFS
- Use Hue visual interface
 - Similar to AWS EMR
- Use Query dropdown menu to select Query>Editor Pig, Hive, Java, Scala, etc
- Use Top-left dropdown menu to select:
 - Apps: Editor, Scheduler
 - Browser: Documents, Files, Tables, Indexes, Jobs
- Use Browsers > Files to Upload your datasets or script
 - Select right button Upload > Files
 - When connected with browser via SSH Tunnel, you will have access to local host directories



Exercises local directory mounted to Cludera FS (typically in the /media share)

The screenshot displays a virtual machine environment with the following components:

- Host OS:** Windows 10 desktop with taskbar and Start menu visible.
- Virtual Machine:** Oracle VM VirtualBox running a Linux-based Cludera environment.
- Cludera Interface:** The Cludera web interface is open in a Mozilla Firefox browser window. It shows a search bar, a sidebar with navigation links (Impala, Spark, Solr, Oozie, Cludera Manager, etc.), and a main content area with a table of jobs.
- File Browser:** A file browser window titled "sf_Exercise_Files_all" is open, showing a directory structure with folders like "Ex_Files_BigData_Hive" and "Ex_Files_Hadoop_Data_Science_Tips".
- Code Editor:** A gedit window titled "7-5.sql (sf_Exercise_Files_all) - gedit" is open, displaying SQL code for creating tables and performing queries.

SQL Code in gedit:


```
-- identify large sales with IF
select
  orderid,
  saleamount,
  if(saleamount > 5000, 1, 0) as LargeSale
from sales_all_years
limit 1000;

-- create sales size categories
select
  orderid,
  saleamount,
  case
    when saleamount > 5000 then 'large'
    when saleamount > 1000 then 'medium'
    else 'small'
  end as SalesSize
from sales_all_years
limit 1000;

-- perform what-if analysis by reassigning regions
select
  case lower(region)
    when 'west' then 'Southwest'
    when 'south' then 'Southwest'
```

Cludera Jobs Table:

Owner	Last Modified
hue	12/28/2019 9:35 AM



Hue: File Browser/Directory and Editor view

File Browser

Search for file name

Actions Delete forever

Upload New

Home / user / demch

	Name	Size	User	Group	Permissions	Date
	↑		hdfs	hadoop	drwxr-xr-x	May 21, 2019 02:18 PM
	.		demch	demch	drwxr-xr-x	May 21, 2019 02:28 PM
	CIS_FacultyList.csv	136.0 KB	demch	demch	-rw-r--r--	May 21, 2019 02:28 PM
		102 bytes	demch	demch	-rw-r--r--	May 21, 2019 02:28 PM
		102 bytes	demch	demch	-rw-r--r--	May 21, 2019 02:28 PM
			demch	demch	drwxr-xr-x	May 21, 2019 02:27 PM

Page 1 of 1

Editor

Dashboard

Scheduler

Impala

Hive

Pig

Java

Spark

MapReduce

Shell

Sqoop 1

Distcp

Sqlr SQL


Add more...

Query History

Results (3)

```
CREATE TABLE students2
(last_name STRING, student_id
BIGINT, course STRING,
start_date DATE, address STRING
COMMENT 'Permanent home address
of the student',
highest_qualification STRING,
degree_type STRING, country
STRING) STORED AS SEQUENCEFILE
```

- When accessing local cluster via browser, you can upload file/data via File Browser in Hue from your local machine



Use File Browser to Upload files

web-digcomp2.1pdf_(online)-intro-tables-examples.pdf - Adobe Acrobat Pro DC

cloudera-quickstart-vm-5.13.0-0-virtualbox [Running] - Oracle VM VirtualBox

File Machine View Input Devices Help

Applications Places System

Wed Apr 15, 7:10 PM cloudera

sf_Exercise_Files_all - File Browser

File Edit View Go Bookmarks Tabs Help

Back Forward

Places

- cloudera
- Desktop
- File System
- Network
- sf_Exercise_...
- Trash
- Documents
- Music
- Pictures
- Videos
- Downloads

Ex_Files_BigData_Hive

Ex_Files_Hadoop_Data_Science_Tips

2 items, Free space: 212.5 GB

```
when saleamount > 5000 then
  when saleamount > 1000 then
    else 'small'
  end as SalesSize
from sales_all_years
limit 1000;

-- perform what-if analysis by region
select
  case lower(region)
    when 'west' then 'Southwest'
    when 'south' then 'South'
```

Hue - File Browser - Mozilla Firefox

Restore Session x Hue - Table Browser x Cloudera Live: Welco... x Hue - Table Browser x Hue - File Browser x Error while proces... x

quickstart.cloudera:8888/hue/filebrowser/view=/user/cloudera

Cloudera Hue Hadoop HBase Impala Spark Solr Oozie Cloudera Manager Getting Started Yuri Demchenko's ...

Query Search data and saved documents...

Jobs cloudera

File Browser

Search for file name Actions Move to trash Upload New

Files Zip/Tgz/Bz2 file

Home / user / cloudera

	Name	Size	User	Group	Permissions	Date
<input type="checkbox"/>	.		hdfs	supergroup	drwxr-xr-x	December 28, 2019 09:32 AM
<input type="checkbox"/>	.		cloudera	cloudera	drwxr-xr-x	December 28, 2019 02:24 PM
<input type="checkbox"/>	2015_11_18		cloudera	cloudera	drwxrwxrwx	December 28, 2019 09:33 AM
<input type="checkbox"/>	2015_11_19		cloudera	cloudera	drwxrwxrwx	December 28, 2019 09:33 AM
<input type="checkbox"/>	2015_11_20		cloudera	cloudera	drwxrwxrwx	December 28, 2019 09:33 AM
<input type="checkbox"/>	2015_11_21		cloudera	cloudera	drwxrwxrwx	December 28, 2019 09:33 AM
<input type="checkbox"/>	CogsleyServices-SalesData-US.xlsx	2.3 MB	cloudera	cloudera	-rw-r--r--	December 28, 2019 02:04 PM
<input type="checkbox"/>	csv-serde-1.1.2-0.11.0-all.jar	24.6 KB	cloudera	cloudera	-rw-r--r--	December 28, 2019 02:24 PM

Show 45 of 6 items Page 1 of 1

Click or Drop files here

Hue - File Browser - M... sf_Exercise_Files_all - ... 7-5.sql (sf_Exercise_Fil...

Right Ctrl

4:14 AM 16-Apr-20

Accessing CDH cluster from your browser

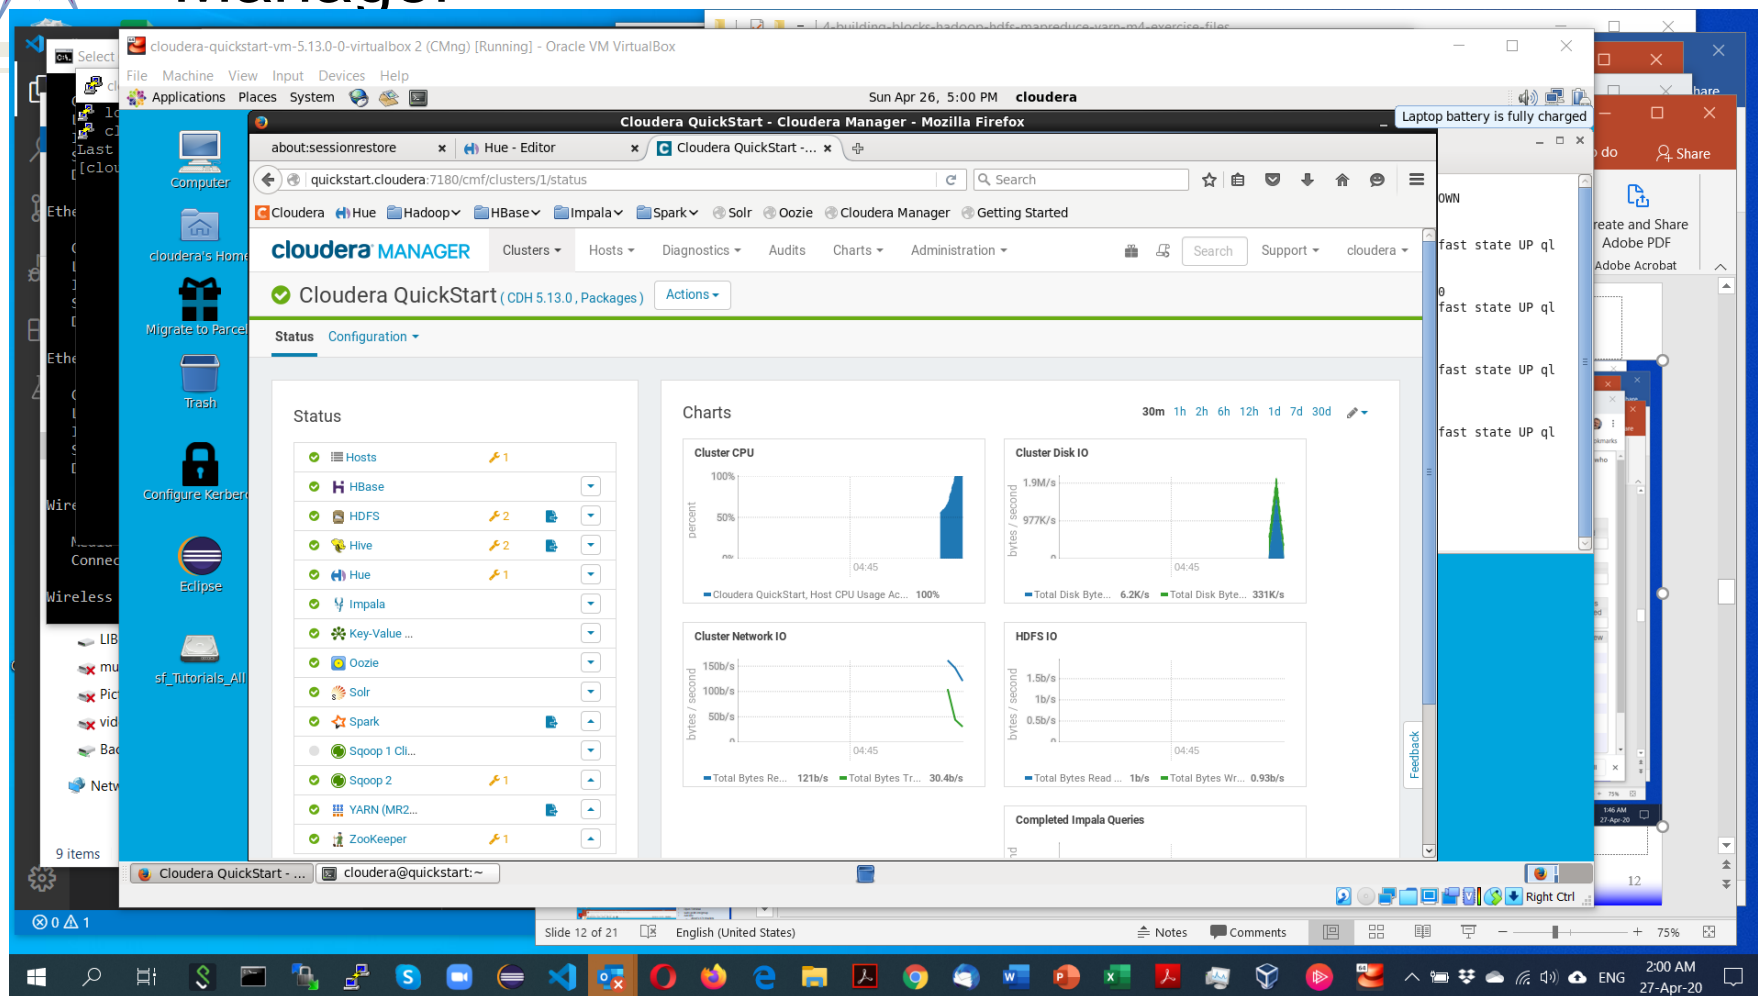
The image displays two overlapping screenshots of the Hue web interface. The top screenshot shows the 'Hue - Editor' window with a SQL query being executed in the Hive database. The query is as follows:

```
1 drop table vip_clients_intl;
2 -- create new vip table
3 create table if not exists vip_clients_intl(name string, country string);
4 -- load some data
5 insert into vip_clients_intl values
6 ('Apple Inc.', 'USA'),
7 ('Google Inc.', 'USA'),
8 ('Facebook, Inc.', 'USA'),
9 ('Amazon.com, Inc.', 'USA'),
10 ('Microsooft', 'USA'),
11 ('SAP', 'Germany');
```


The bottom screenshot shows the 'Hue - File Browser' window. An 'Upload to /user/cloudera' dialog box is open, and a file named 'CIS_FacultyList.csv' is being selected from the local file system. The file browser shows a list of files and folders, including 'J', '.Trash', and various date-based folders.

- Use SSH tunnel with convenient port mapping
- **Benefit from accessing local file system and uploading files and data via browser**

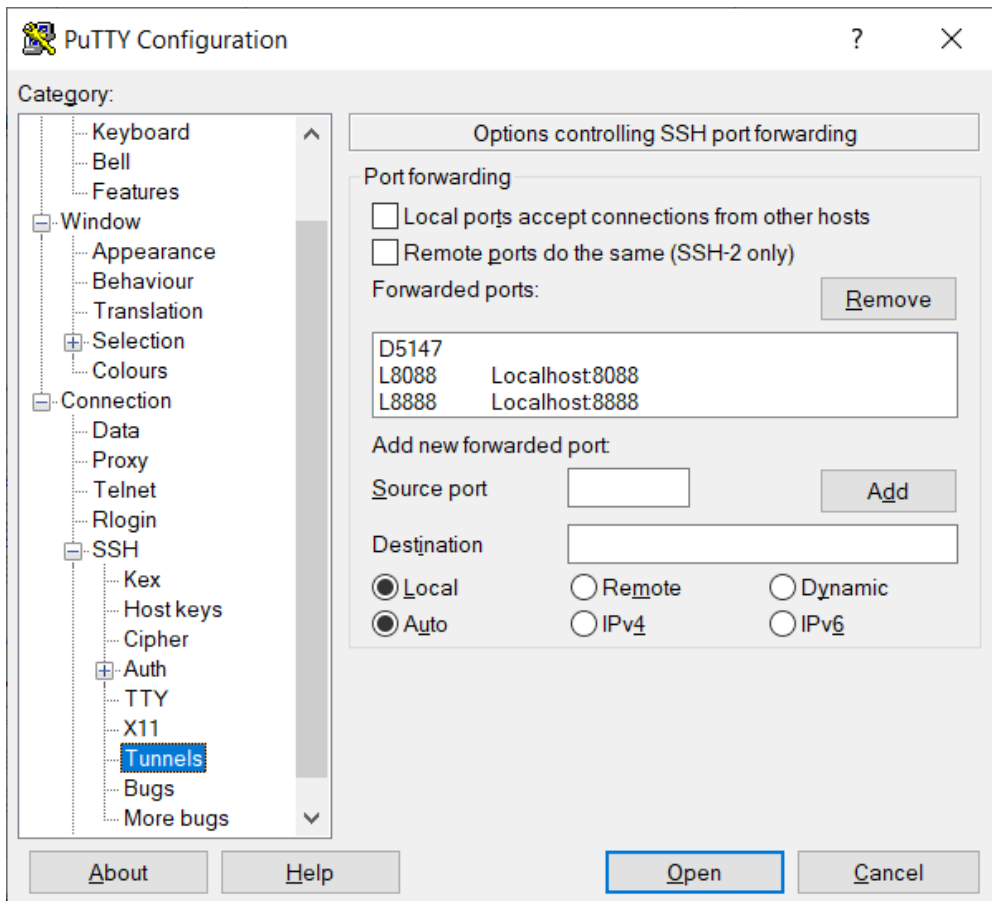
Cloudera Quickstart with enabled Cloudera Manager




- To run Cloudera Manager Launch Cloudera Express from Desktop
- Be aware about requirement for configuring at least 2 virtual CPU and minimum memory 8 GB



Configuring Tunnels for known Ports




- Using PuTTY SSH client
- First, configure dynamic tunnel to EMR cluster
 - Port 8157
 - Using Source port and Dynamic radio button
 - Example:
 - D5147
 - Source L8088
Destination Localhost:8088
 - Source L8888
Destination Localhost:8888
- Configure other local ports
 - See EMR ports table
- This will create a secure tunnel by forwarding a port (the “destination port”) on the remote server to a port (the “source port”) on the local host (127.0.0.1 or localhost).



Configuring SSH tunnel on Windows

<https://docs.bitnami.com/virtual-machine/faq/get-started/access-phpmyadmin/>

- In order to access phpMyAdmin via SSH tunnel, you need an SSH client. In the instructions below we have selected PuTTY, a free SSH client for Windows and UNIX platforms.
 - The first step is to configure PuTTY.
- Once you have your SSH client correctly configured and you have confirmed that you can successfully access your instance using SSH, you need to create an SSH tunnel in order to access phpMyAdmin. Follow these steps:
 - In the “Connection -> SSH -> Tunnels” section, add a new forwarded port by introducing the following values:
 - Source port: 8888
 - Destination: localhost:8888
 - Remember that if you are redirecting HTTP requests to the HTTPS port, you must use destination port 443 instead of 80.
 - This will create a secure tunnel by forwarding a port (the “destination port”) on the remote server to a port (the “source port”) on the local host (127.0.0.1 or localhost).
 - Click the “Add” button to add the secure tunnel configuration to the session. You’ll see the added port in the list of “Forwarded ports”.
- PuTTY configuration
 - In the “Session” section, save your changes by clicking the “Save” button.
 - Click the “Open” button to open an SSH session to the server. The SSH session will now include a secure SSH tunnel between the two specified ports.
- Access the phpMyAdmin console through the secure SSH tunnel you created, by browsing to `http://127.0.0.1:8888/phpmyadmin`.
- Log in to phpMyAdmin by using the following credentials:
 - Username: root
 - Password: application password. (Refer to our FAQ to learn how to find your application credentials).




Running MapReduce examples on CDH5.13

Follow step, adjust jar name and commands

https://docs.cloudera.com/documentation/other/tutorial/CDH5/topics/ht_usage.html#topic_5_2


- `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar wordcount /user/cloudera/wordcount/output`
- `$ hadoop fs -rm -r /user/cloudera/wordcount/output`
- `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 10 10000`
- `$ hadoop fs -rm -r /user/cloudera/wordcount/output`
- `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar wordmean /user/cloudera/wordcount/output`



Commands/Programs in hadoop-mapreduce-examples.jar

Jobs in the examples JAR file hadoop-mapreduce-examples.jar

- **aggregatewordcount**: An Aggregate-based map/reduce program that counts the words in the input files.
- **aggregatewordhist**: An Aggregate-based map/reduce program that computes the histogram of the words in the input files.
- **bbp**: A map/reduce program that uses Bailey-Borwein-Plouffe to compute the exact digits of pi.
- **dbcoun**: An example job that counts the pageview counts from a database.
- **distbbp**: A map/reduce program that uses a BBP-type formula to compute the exact bits of pi.
- **grep**: A map/reduce program that counts the matches to a regex in the input.
- **join**: A job that effects a join over sorted, equally partitioned data sets.
- **multifilewc**: A job that counts words from several files.
- **pentomino**: A map/reduce tile laying program to find solutions to pentomino problems.
- **pi**: A map/reduce program that estimates pi using a quasi-Monte Carlo method.
- **randomtextwriter**: A map/reduce program that writes 10 GB of random textual data per node.
- **randomwriter**: A map/reduce program that writes 10 GB of random data per node.
- **secondarysort**: An example defining a secondary sort to the reduce.
- **sort**: A map/reduce program that sorts the data written by the random writer.
- **sudoku**: A Sudoku solver.
- **teragen**: Generate data for the terasort.
- **terasort**: Run the terasort.
- **teravalidate**: Check the results of the terasort.
- **wordcount**: A map/reduce program that counts the words in the input files.
- **wordmean**: A map/reduce program that counts the average length of the words in the input files.
- **wordmedian**: A map/reduce program that counts the median length of the words in the input files.
- **wordstandarddeviation**: A map/reduce program that counts the standard deviation of the length of the words in the input files.




Do you need Reduce?

ID	Username	Category	Amount
1	Janani	Books	200
2	Swetha	Clothing	450
3	Shreya	Electronics	300
4	Jitu	Books	700

Which users spent >300?

**How many of them
bought Books?**

**Selecting a specific set
of records from a dataset**



SQL Query can do it


Selecting a specific set of records

If this were a database table

An SQL query

```
select * from <table name>
where <condition>
```

ID	Username	Category	Amount
1	Janani	Books	200
2	Swetha	Clothing	450
3	Shreya	Electronics	300
4	Jitu	Books	700



Run MapReduce for Standard Select Operation

Get the top N records

If this were a database table

An SQL query
order by

```
select * from <table name>  
where <condition>  
order by <column name>
```

User	Followers
1	30
2	30000
3	20
4	40
5	50
6	6000